## Lecture outline

- Support vector machines
- Boosting


## Support Vector Machines



Find a linear hyperplane (decision boundary) that will separate the data

## Support Vector Machines



- One Possible Solution


## Support Vector Machines



- Another possible solution


## Support Vector Machines



- Other possible solutions


## Support Vector Machines



- Which one is better? B1 or B2?
- How do you define better?


## Support Vector Machines <br> 

- Find hyperplane maximizes the margin $=>$ B1 is better than B2


## Support Vector Machines



## Support Vector Machines

- We want to maximize: $\operatorname{Margin}=\frac{2}{\|w\|^{2}}$
- Which is equivalent to minimizing: $L(w)=\frac{\|w\|^{2}}{2}$
- But subjected to the following constraints:

$$
f\left(\vec{x}_{i}\right)=\left\{\begin{array}{cc}
1 & \text { if } \overrightarrow{\mathrm{w}} \bullet \overrightarrow{\mathrm{x}}_{\mathrm{i}}+\mathrm{b} \geq 1 \\
-1 & \text { if } \overrightarrow{\mathrm{w}} \bullet \overrightarrow{\mathrm{x}}_{\mathrm{i}}+\mathrm{b} \leq-1
\end{array}\right.
$$

- This is a constrained optimization problem
- Numerical approaches to solve it (e.g., quadratic programming)


## Support Vector Machines

- What if the problem is not linearly separable?



## Support Vector Machines

- What if the problem is not linearly separable?
- Introduce slack variables
- Need to minimize:
- Subject to:

$$
L(w)=\frac{\|\vec{w}\|^{2}}{2}+C\left(\sum_{i=1}^{N} \xi_{i}^{k}\right)
$$

$$
f\left(\vec{x}_{i}\right)=\left\{\begin{array}{cc}
1 & \text { if } \overrightarrow{\mathrm{w}} \bullet \overrightarrow{\mathrm{x}}_{\mathrm{i}}+\mathrm{b} \geq 1-\xi_{\mathrm{i}} \\
-1 & \text { if } \overrightarrow{\mathrm{w}} \bullet \overrightarrow{\mathrm{x}}_{\mathrm{i}}+\mathrm{b} \leq-1+\xi_{\mathrm{i}}
\end{array}\right.
$$

## Nonlinear Support Vector Machines

- What if decision boundary is not linear?



## Nonlinear Support Vector Machines

- Transform data into higher dimensional space



## Ensemble Methods

- Construct a set of classifiers from the training data
- Predict class label of previously unseen records by aggregating predictions made by multiple classifiers


## Why does it work?

- Suppose there are 25 base classifiers
- Each classifier has error rate, $\varepsilon=0.35$
- Assume classifiers are independent
- Probability that the ensemble classifier makes a wrong prediction:

$$
\sum_{i=13}^{25} \epsilon^{i}(1-\epsilon)^{25-i}=0.06
$$

## Examples of Ensemble Methods

- How to generate an ensemble of classifiers?
- Bagging
- Boosting


## Bagging

- Sampling with replacement

| Original Data | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
| :--- | :---: | :---: | :---: | :---: | :---: | :---: | :---: | :---: | :---: | :---: |
| Bagging (Round 1) | 7 | 8 | 10 | 8 | 2 | 5 | 10 | 10 | 5 | 9 |
| Bagging (Round 2) | 1 | 4 | 9 | 1 | 2 | 3 | 2 | 7 | 3 | 2 |
| Bagging (Round 3) | 1 | 8 | 5 | 10 | 5 | 5 | 9 | 6 | 3 | 7 |

- Build classifier on each bootstrap sample
- Each sample has probability $1-(1-1 / n)^{n}$ of being selected


## Boosting

- An iterative procedure to adaptively change distribution of training data by focusing more on previously misclassified records
- Initially, all N records are assigned equal weights
- Unlike bagging, weights may change at the end of boosting round


## Boosting

- Records that are wrongly classified will have their weights increased
- Records that are classified correctly will have their weights decreased

| Original Data | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 |
| :--- | :---: | :---: | :---: | :---: | :---: | :---: | :---: | :---: | :---: | :---: |
| Boosting (Round 1) | 7 | 3 | 2 | 8 | 7 | 9 | 4 | 10 | 6 | 3 |
| Boosting (Round 2) | 5 | 4 | 9 | 4 | 2 | 5 | 1 | 7 | 4 | 2 |
| Boosting (Round 3) | 4 | 4 | 8 | 10 | 4 | 5 | 4 | 6 | 3 | 4 |

- Example 4 is hard to classify
- Its weight is increased, therefore it is more likely to be chosen again in subsequent rounds


## Example: AdaBoost

- Base classifiers: $\mathrm{C}_{1}, \mathrm{C}_{2}, \ldots, \mathrm{C}_{\mathrm{T}}$
- Data pairs: $\left(x_{i}, y_{i}\right)$
- Error rate:
$\epsilon_{i}=\frac{1}{N} \sum_{j=1}^{N} w_{j} \delta\left(C_{i}\left(x_{j}\right) \neq y_{j}\right)$
- Importance of a classifier:

$$
\alpha_{i}=\frac{1}{2} \log \left(\frac{1-\epsilon_{i}}{\epsilon_{i}}\right)
$$



## Example: AdaBoost

- Classification:

$$
C^{*}=\arg \max _{y} \sum_{j=1}^{T} \alpha_{j} \delta\left(C_{j}(x)=y\right)
$$

- Weight update for every iteration $t$ and classifier j :

$$
w_{i}^{(t+1)}=\frac{w_{i}^{(t)}}{Z_{t}} \begin{cases}\exp ^{-\alpha_{j}} & \text { if } C_{j}\left(x_{i}\right)=y_{i} \\ \exp ^{\alpha_{j}} & \text { if } C_{j}\left(x_{i}\right) \neq y_{i}\end{cases}
$$

where $Z_{j}$ is the normalization factor

- If any intermediate rounds produce error rate higher than $50 \%$, the weights are reverted back to $1 / n$

